

Attrition prediction

Project 4

Members:

Dale Munroe

Nick Taheri

Jose Sandoval

Overview

Dataset

Data Source:

<https://www.kaggle.com/datasets/pavansubhasht/ibm-hr-analytics-attrition-dataset>

Original Source Format: *csv “imported into postgresql”*

Dependent Variable : *Attrition* - (Boolean Y/N)

Independent Variable(s): 35 (e.g. basic info, work Info, satisfaction, salary and time related)

Data exploration & cleaning.

Data display & null values check

```
#Create DataFrame from SQL table
pg_connection = f'{protocol}://{username}:{password}@{host}:{port}/{database_name}'
conn = pg.connect(pg_connection)
raw_ibm_df = psql.read_sql('SELECT * FROM ibm_employee_data', conn)

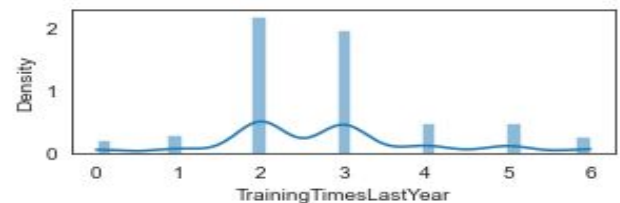
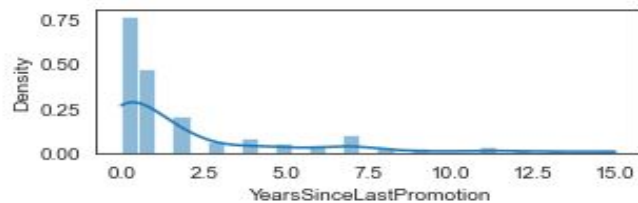
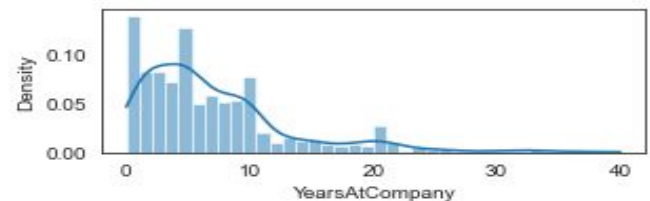
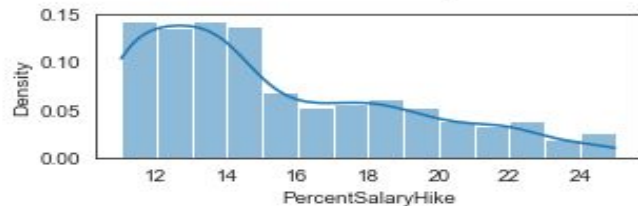
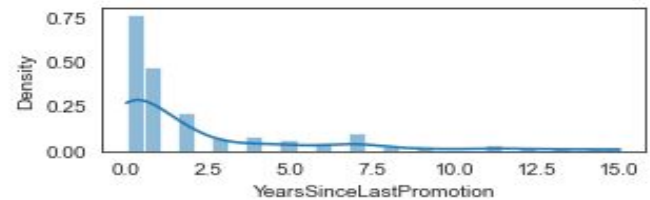
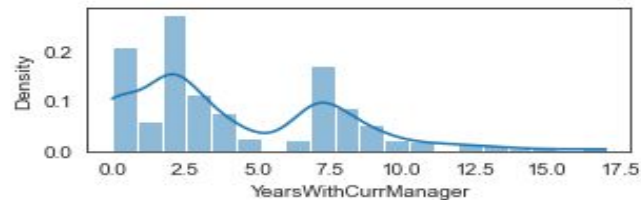
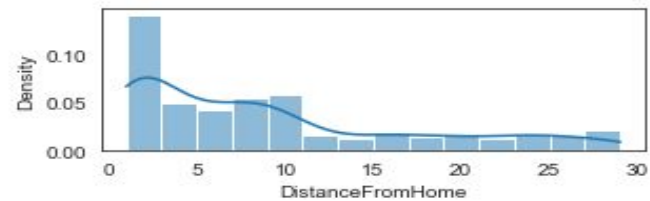
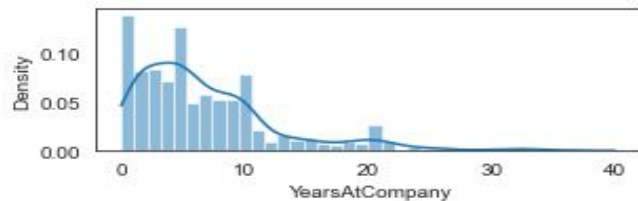
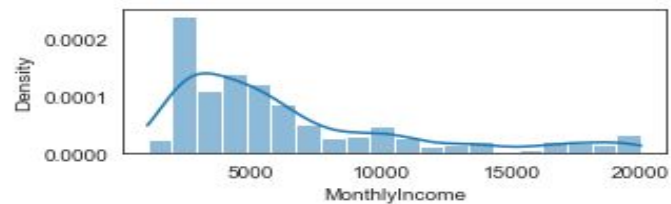
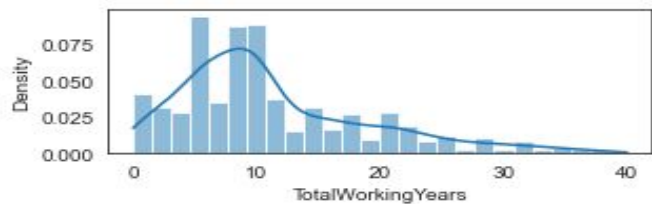
# Display all the DataFrame Columns
pd.options.display.max_columns = None
display(raw_ibm_df)
```

```
# check missing data for each feature
df[df.isnull().any(axis=1)]
```

✓ 0.3s

	age	attrition	businesstravel	dailyrate	department	distancefromhome	education	educationfield	employeecount	employeenumber	environmentsatisfaction	gender	hourlyrate	jobinvolvement	joblevel
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	1	2	Female	94	3	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	2	3	Male	61	2	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	4	4	Male	92	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	5	4	Female	56	3	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	7	1	Male	40	3	

Data Distribution



Insignificant Data

Columns with only 1 unique value & the employee # was removed

Age	43
Attrition	2
BusinessTravel	3
DailyRate	886
Department	3
DistanceFromHome	29
Education	5
EducationField	6
EmployeeCount	1
EmployeeNumber	1470
EnvironmentSatisfaction	4
Gender	2
HourlyRate	71
JobInvolvement	4
JobLevel	5
JobRole	9
JobSatisfaction	4
MaritalStatus	3
MonthlyIncome	1349
MonthlyRate	1427
NumCompaniesWorked	10
Over18	1
Overtime	2
PercentSalaryHike	15
PerformanceRating	2
RelationshipSatisfaction	4
StandardHours	1
StockOptionLevel	4
TotalWorkingYears	40
TrainingTimesLastYear	7
WorkLifeBalance	4
YearsAtCompany	37
YearsInCurrentRole	19
YearsSinceLastPromotion	16
YearsWithCurrManager	18
dtype: int64	

```
# drop columns that are irrelevant
df = df.drop(['Over18', 'EmployeeNumber', 'EmployeeCount', 'StandardHours'], axis = 1)
print(df.shape)
df.head()
```

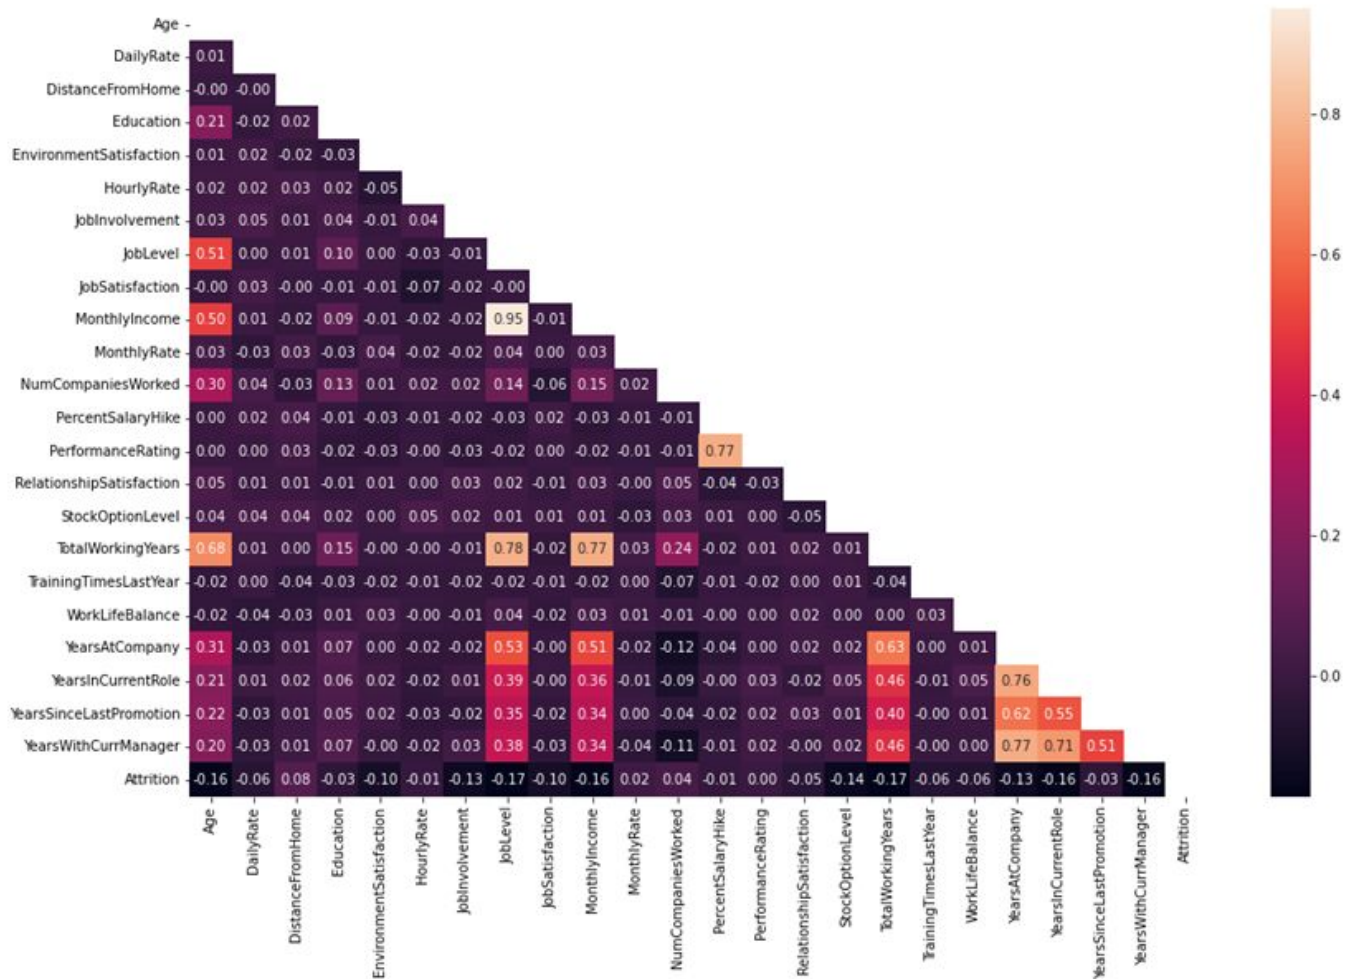
✓ 0.7s

(1470, 31)

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EnvironmentSatisfaction	Gender	HourlyRate	JobInvolvement
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	2	Female	94	3
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	3	Male	61	2
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	4	Male	92	2
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	4	Female	56	3
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	Male	40	3

Feature Engineering

- Correlation Matrix for variables
- One Hot Encoding



```
# apply one hot encoding to non numerical features
df_clean = pd.get_dummies(df_clean, columns = ['BusinessTravel', 'Gender', 'MaritalStatus'], drop_first = True)
df_clean = pd.get_dummies(df_clean)
df_clean.head()
```

✓ 0.1s

	Age	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	HourlyRate	JobInvolvement	JobLevel	JobSatisfaction	MonthlyIncome	MonthlyRate	NumCompaniesWorked	OverTime	PercentSalaryHike
0	41	1102	1	2	2	94	3	2	4	5993	19479	8	1	11
1	49	279	8	1	3	61	2	2	2	5130	24907	1	0	23
2	37	1373	2	2	4	92	2	1	3	2090	2396	6	1	15
3	33	1392	3	4	4	56	3	1	3	2909	23159	1	1	11
4	27	591	2	1	1	40	3	1	2	3468	16632	9	0	12

Data Scaling

```
# filter out features that needs to be standardized
col_tobe_standard = ['Age', 'DailyRate', 'DistanceFromHome', 'Education', 'EnvironmentSatisfaction',
                     'HourlyRate', 'JobInvolvement', 'JobLevel', 'JobSatisfaction', 'MonthlyIncome',
                     'MonthlyRate', 'NumCompaniesWorked', 'PercentSalaryHike', 'PerformanceRating',
                     'RelationshipSatisfaction', 'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
                     'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',
                     'YearsWithCurrManager']
```

✓ 0.5s

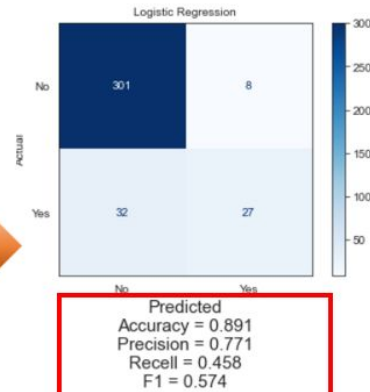
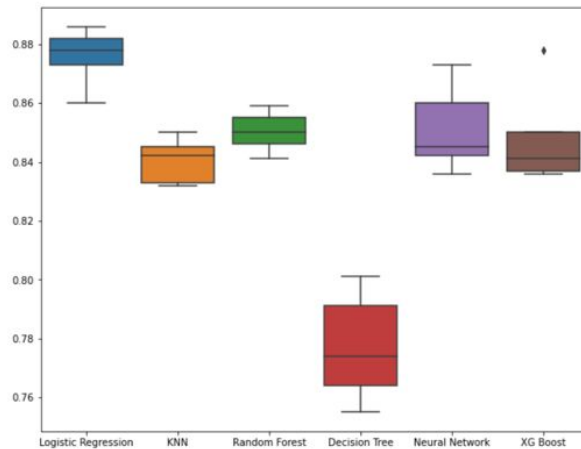
```
# standarization on numercial features so that all the numerical features are having the same type of normal distribution
scaler = StandardScaler()
for col in col_tobe_standard:
    df_clean[col] = df_clean[col].astype(float)
    df_clean[[col]] = scaler.fit_transform(df_clean[[col]])
df_clean.head()
```

✓ 0.3s

	Age	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	HourlyRate	JobInvolvement	JobLevel	JobSatisfacti
0	0.446350	0.742527	-1.010909	-0.891688	-0.660531	1.383138	0.379672	-0.057788	1.1532
1	1.322365	-1.297775	-0.147150	-1.868426	0.254625	-0.240677	-1.026167	-0.057788	-0.6608
2	0.008343	1.414363	-0.887515	-0.891688	1.169781	1.284725	-1.026167	-0.961486	0.2462
3	-0.429664	1.461466	-0.764121	1.061787	1.169781	-0.486709	0.379672	-0.961486	0.2462
4	-1.086676	-0.524295	-0.887515	-1.868426	-1.575686	-1.274014	0.379672	-0.961486	-0.6608

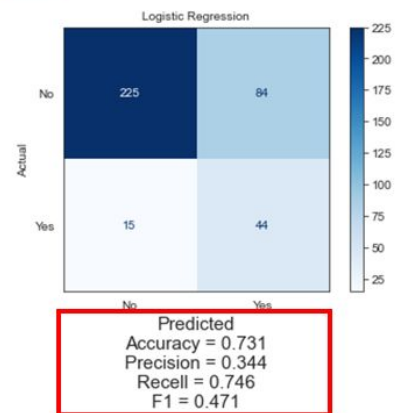
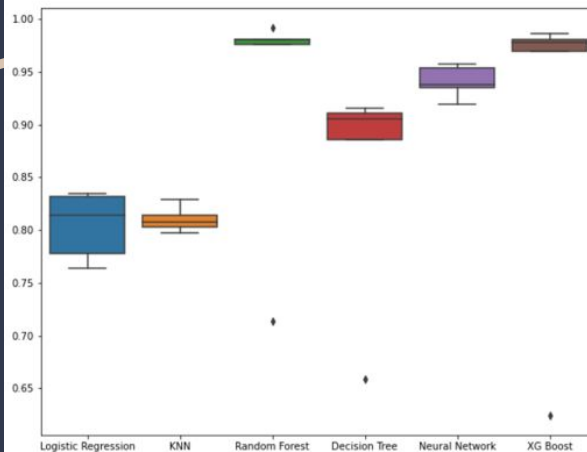
89 %
To good to be true?
Do we need to deal
with data imbalance /
overprediction

Data Imbalance Not Addressed



"To good to be true?"

Data Imbalance Not Addressed (SMOTE)



Challenges Data Imbalance

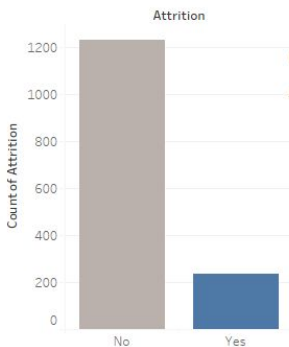
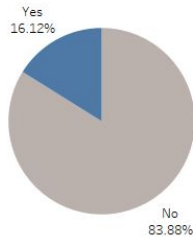
Solution

imblearn

- Oversampling
- Undersampling
- SMOTE

(Synthetic Minority Oversampling Technique)

IBM



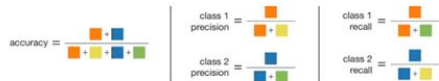
Q: How do we Deal with a Imbalanced DataSet
A: Use imblearn Library



	Predicted label class 1	Predicted label class 2
True label class 1	correct true positive for class 1	wrong false positive for class 2
True label class 2	wrong false positive for class 1	correct true positive for class 2

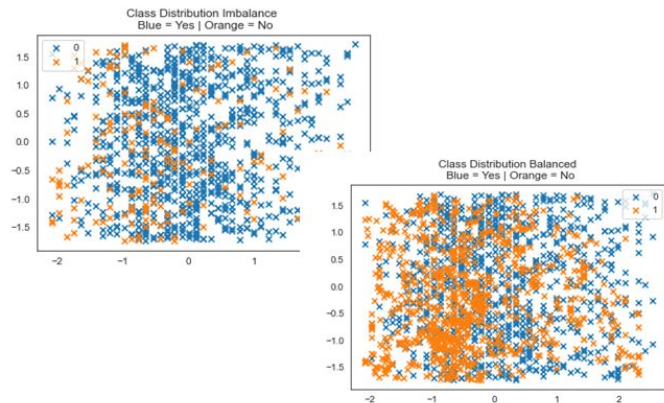
Using: Oversampling & Undersampling

Using SMOTE



“ACCURACY” is not the right matrix when working for the Imbalanced data set. When you use the confusion matrix, we can evaluate how well the model is really performing.

- The accuracy of the model is basically the total number of correct predictions divided by total number of predictions.
- The precision of a class define how trustable is the result when the model answer that a point belongs to that class.
- The recall of a class expresses how well the model is able to detect that class.
- The F1 score of a class is given by the harmonic mean of precision and recall ($2 \times \text{precision} \times \text{recall} / (\text{precision} + \text{recall})$), it combines precision and recall of a class in one metric.



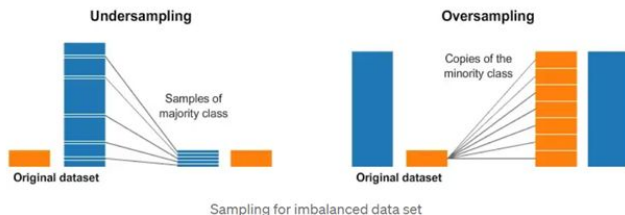
Oversampling & Undersampling



Using Under & Over Sampling Methods

"ACCURACY" is not the right matrix when working for the Imbalanced data set. When you use the confusion matrix, we can evaluate how well the model is really performing.

- The accuracy of the model is basically the total number of correct predictions divided by total number of predictions.
- The precision of a class define how trustable is the result when the model answer that a point belongs to that class.
- The recall of a class expresses how well the model is able to detect that class.
- The F1 score of a class is given by the harmonic mean of precision and recall ($2 \times \text{precision} \times \text{recall} / (\text{precision} + \text{recall})$), it combines precision and recall of a class in one metric.



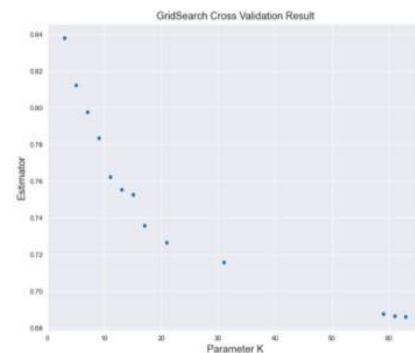
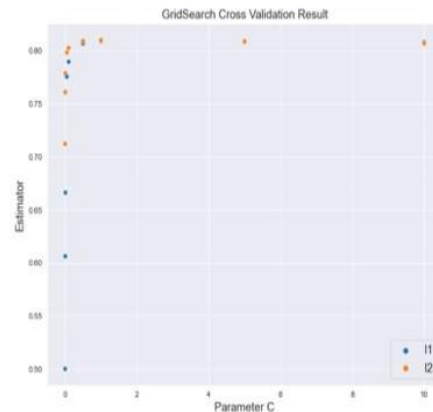
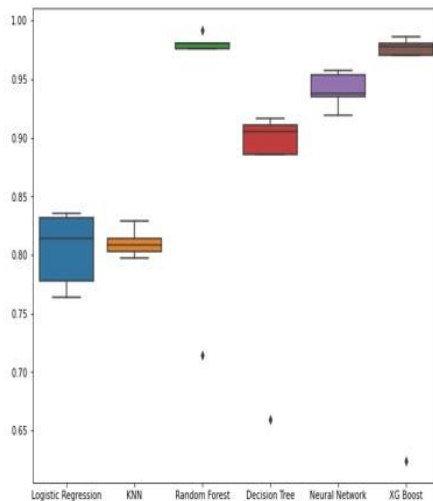
	method	mean score	std score	recall score
0	Logistic Regression	0.315079	0.315079	0.322034
1	KNN Classifier	0.140794	0.032205	0.169492
2	Random Forest Classifier	0.106984	0.060737	0.220339
3	Decision Tree Classifier	0.400159	0.100461	0.474576
4	MLP Classifier	0.196984	0.095284	0.474576
5	XG Boost	0.297937	0.043541	0.322034

	method	mean score	std score	recall score
0	Logistic Regression UnderSampling	0.315079	0.702063	0.779661
1	KNN Classifier UnderSampling	0.568730	0.127682	0.762712
2	Random Forest Classifier UnderSampling	0.697460	0.079316	0.762712
3	Decision Tree Classifier UnderSampling	0.674127	0.094545	0.559322
4	MLP Classifier UnderSampling	0.592381	0.205470	0.966102
5	XG Boost UnderSampling	0.714127	0.087748	0.728814

	method	mean score	std score	recall score
0	Logistic Regression OverSampling	0.702063	0.063689	0.796610
1	KNN Classifier OverSampling	0.568730	0.127682	0.762712
2	Random Forest Classifier OverSampling	0.202222	0.045051	0.338983
3	Decision Tree Classifier OverSampling	0.263651	0.073520	0.491525
4	MLP Classifier OverSampling	0.752540	0.180019	0.949153
5	XG Boost OverSampling	0.330952	0.059370	0.474576

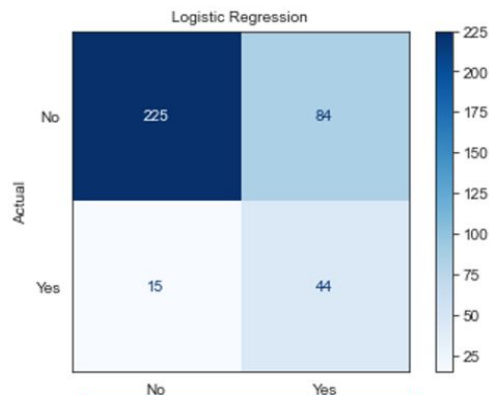
SMOTE

Model Training and Performance Evaluation

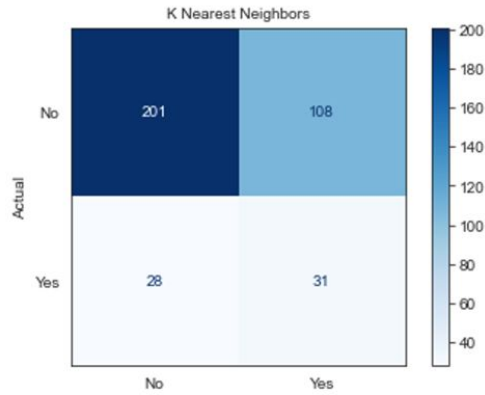


Confusion Matrix

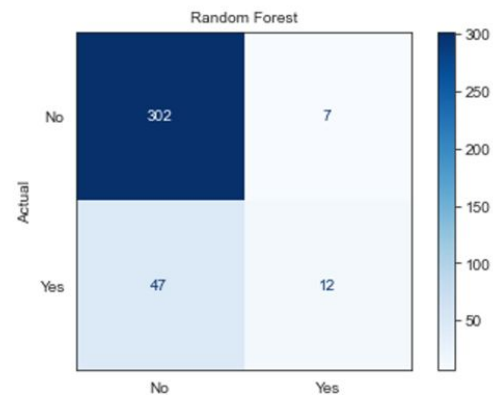
SMOTE



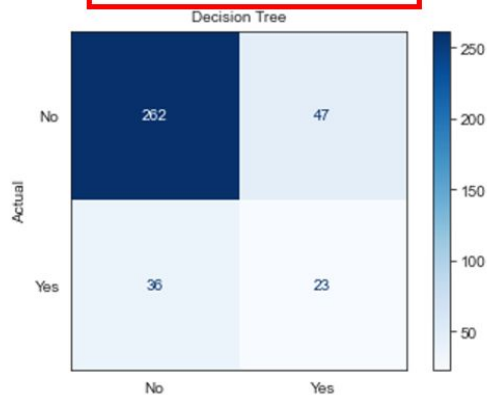
Predicted
Accuracy = 0.731
Precision = 0.344
Recall = 0.746
F1 = 0.471



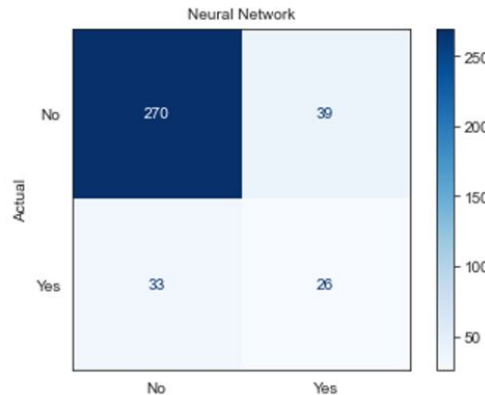
Predicted
Accuracy = 0.63
Precision = 0.223
Recall = 0.525
F1 = 0.313



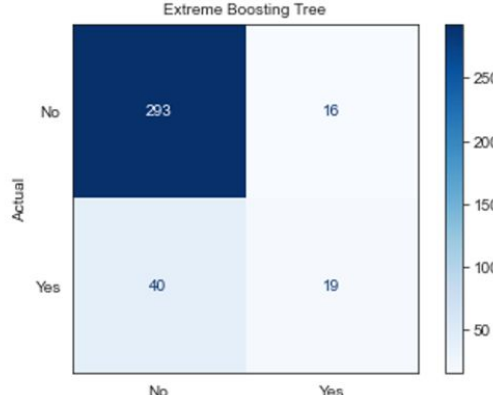
Predicted
Accuracy = 0.853
Precision = 0.632
Recall = 0.203
F1 = 0.308



Predicted
Accuracy = 0.774
Precision = 0.329
Recall = 0.39
F1 = 0.357



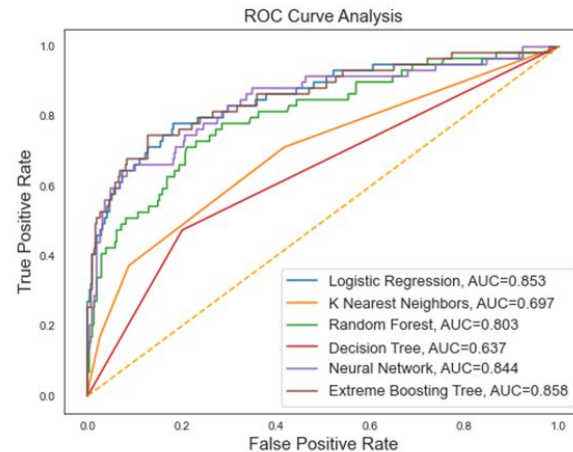
Predicted
Accuracy = 0.804
Precision = 0.4
Recall = 0.441
F1 = 0.419



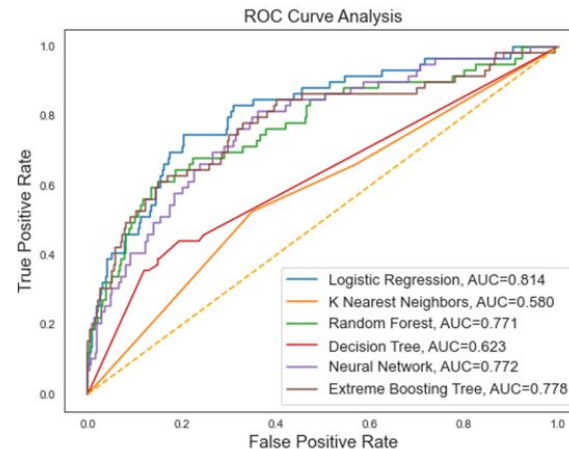
Predicted
Accuracy = 0.848
Precision = 0.543
Recall = 0.322
F1 = 0.404

ROC Curve Analysis Comparison

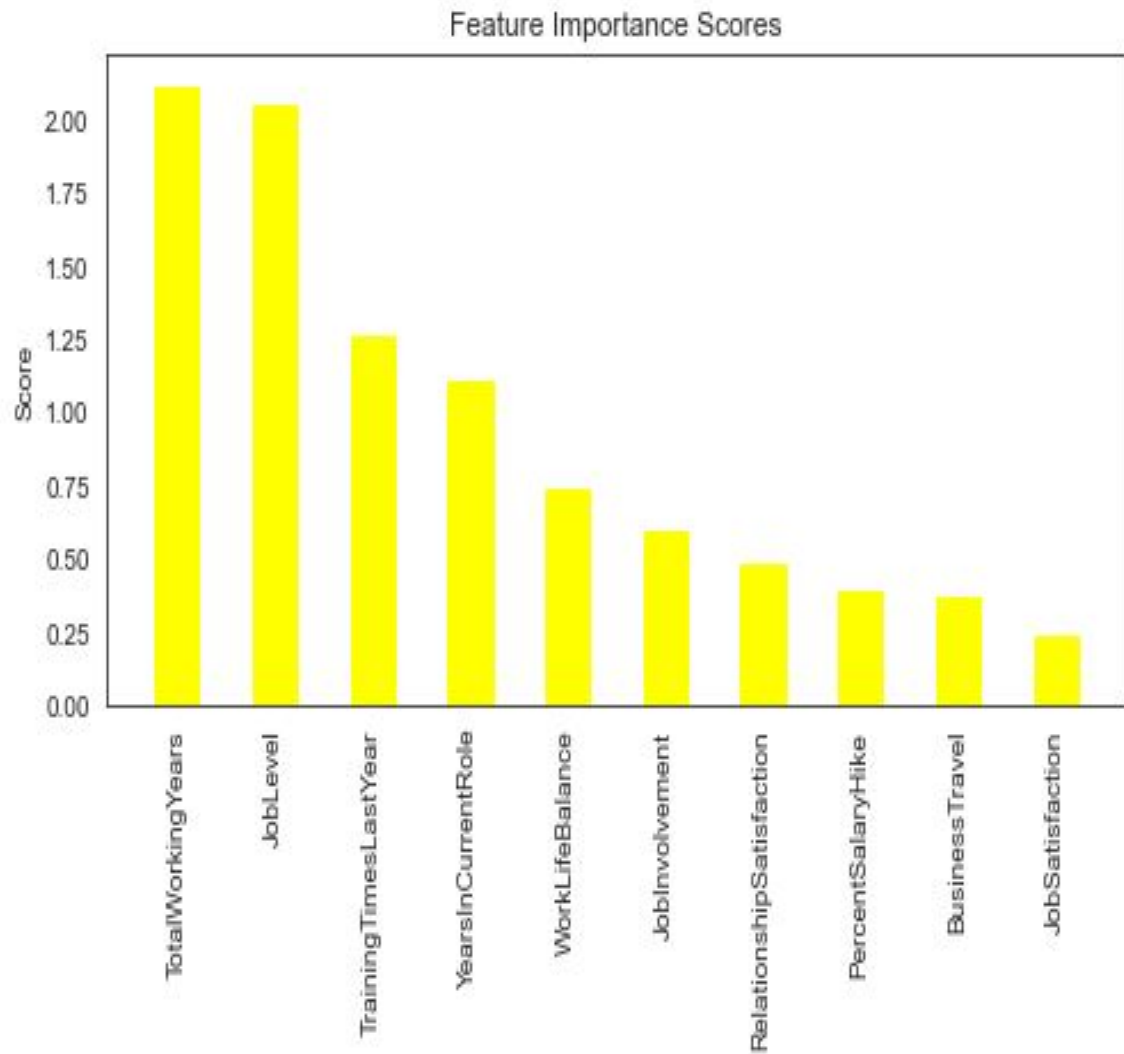
Data Imbalance Not Addressed



Data Imbalance Not Addressed (SMOTE)



Feature Importance



Conclusions

“Thankyou
<https://medium.com/>”

Conclusions from the Project.

We built a model that was achieving a score between 77 – 88 % depending on the method used.

On reflection this is to be expected as the data is bias towards NO. (84% to 16%).

So, the model learns that if it picks NO, it will be right 84% of the time. The model was doing the right thing but the way of training the model is wrong and we are focusing on the wrong thing.

“ACCURACY” is not the right matrix when working for the Imbalanced data set.

When you use the confusion matrix, we can evaluate how well the model is really performing.

		Predicted label class 1	Predicted label class 2
True label class 1	<div>correct true positive for class 1</div>	<div>wrong false positive for class 2</div>	
True label class 2	<div>wrong false positive for class 1</div>	<div>correct true positive for class 2</div>	

$$\text{accuracy} = \frac{\text{orange} + \text{blue}}{\text{orange} + \text{yellow} + \text{blue} + \text{green}}$$
$$\begin{array}{l} \text{class 1 precision} = \frac{\text{orange}}{\text{orange} + \text{yellow}} \\ \text{class 2 precision} = \frac{\text{blue}}{\text{blue} + \text{green}} \end{array}$$
$$\begin{array}{l} \text{class 1 recall} = \frac{\text{orange}}{\text{orange} + \text{green}} \\ \text{class 2 recall} = \frac{\text{blue}}{\text{blue} + \text{yellow}} \end{array}$$

- The accuracy of the model is basically the total number of correct predictions divided by total number of predictions.
- The precision of a class define how trustable is the result when the model answer that a point belongs to that class.
- The recall of a class expresses how well the model can detect that class.
- The F1 score of a class is given by the harmonic mean of precision and recall ($2 \times \text{precision} \times \text{recall} / (\text{precision} + \text{recall})$), it combines precision and recall of a class in one metric.